

# Exercise

For this exercise, you will be working with the [House Price Dataset](https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques/code) (<https://www.kaggle.com/competitions/house-prices-advanced-regression-techniques/code>).

Please grab the train.csv file from Kaggle and explore this dataset. You need to perform exploratory data analysis and see if there is any correlation between the variables and analyze the distribution of the dataset. The question is open-ended and basically you're asked to perform EDA.

1- Write a summary of your findings in one page (e.g., summary statistics, plots) and submit the pdf file. Therefore, for part 3 of your assignment, you need to submit at least one jupyter notebook file and one pdf file.

2- Push your code and project to github and provide the link to your code here. Ensure that your github project is organized to at least couple of main folders, ensure that you have the README file as well:

- Src
- Data
- Docs
- Results

Read this link for further info: <https://gist.github.com/ericmjl/27e50331f24db3e8f957d1fe7bbbe510> (<https://gist.github.com/ericmjl/27e50331f24db3e8f957d1fe7bbbe510>).

In [15]:

```
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
```


In [16]:

```
df = pd.read_csv("train.csv")
df
```

Out[16]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandConto
0	1	60	RL	65.0	8450	Pave	NaN	Reg	l
1	2	20	RL	80.0	9600	Pave	NaN	Reg	l
2	3	60	RL	68.0	11250	Pave	NaN	IR1	l
3	4	70	RL	60.0	9550	Pave	NaN	IR1	l
4	5	60	RL	84.0	14260	Pave	NaN	IR1	l
...	...	...	...	...	...	...	...	...	...
1455	1456	60	RL	62.0	7917	Pave	NaN	Reg	l
1456	1457	20	RL	85.0	13175	Pave	NaN	Reg	l
1457	1458	70	RL	66.0	9042	Pave	NaN	Reg	l
1458	1459	20	RL	68.0	9717	Pave	NaN	Reg	l
1459	1460	20	RL	75.0	9937	Pave	NaN	Reg	l

1460 rows × 81 columns



## Exploratory Data Analysis

We can see that the dataset consists of 38 numerical columns of 81 all. Moreover, some columns contain missing data.

In [27]:

```
print(df.shape)
```

(1460, 81)

In [28]:

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                    1460 non-null   int64
1   MSSubClass            1460 non-null   int64
2   MSZoning              1460 non-null   object
3   LotFrontage          1201 non-null   float64
4   LotArea              1460 non-null   int64
5   Street               1460 non-null   object
6   Alley               91 non-null     object
7   LotShape             1460 non-null   object
8   LandContour          1460 non-null   object
9   Utilities            1460 non-null   object
10  LotConfig            1460 non-null   object
11  LandSlope            1460 non-null   object
12  Neighborhood         1460 non-null   object
13  Condition1           1460 non-null   object
14  Condition2           1460 non-null   object
15  BldgType             1460 non-null   object
16  HouseStyle           1460 non-null   object
17  OverallQual          1460 non-null   int64
18  OverallCond          1460 non-null   int64
19  YearBuilt            1460 non-null   int64
20  YearRemodAdd         1460 non-null   int64
21  RoofStyle            1460 non-null   object
22  RoofMatl            1460 non-null   object
23  Exterior1st         1460 non-null   object
24  Exterior2nd         1460 non-null   object
25  MasVnrType          1452 non-null   object
26  MasVnrArea          1452 non-null   float64
27  ExterQual            1460 non-null   object
28  ExterCond            1460 non-null   object
29  Foundation          1460 non-null   object
30  BsmtQual            1423 non-null   object
31  BsmtCond            1423 non-null   object
32  BsmtExposure        1422 non-null   object
33  BsmtFinType1        1423 non-null   object
34  BsmtFinSF1          1460 non-null   int64
35  BsmtFinType2        1422 non-null   object
36  BsmtFinSF2          1460 non-null   int64
37  BsmtUnfSF           1460 non-null   int64
38  TotalBsmtSF         1460 non-null   int64
39  Heating             1460 non-null   object
40  HeatingQC           1460 non-null   object
41  CentralAir          1460 non-null   object
42  Electrical          1459 non-null   object
43  1stFlrSF            1460 non-null   int64
44  2ndFlrSF            1460 non-null   int64
45  LowQualFinSF        1460 non-null   int64
46  GrLivArea           1460 non-null   int64
47  BsmtFullBath        1460 non-null   int64
48  BsmtHalfBath        1460 non-null   int64
49  FullBath            1460 non-null   int64
50  HalfBath            1460 non-null   int64
```

```

51 BedroomAbvGr    1460 non-null    int64
52 KitchenAbvGr    1460 non-null    int64
53 KitchenQual      1460 non-null    object
54 TotRmsAbvGrd    1460 non-null    int64
55 Functional       1460 non-null    object
56 Fireplaces       1460 non-null    int64
57 FireplaceQu      770 non-null     object
58 GarageType       1379 non-null    object
59 GarageYrBlt      1379 non-null    float64
60 GarageFinish     1379 non-null    object
61 GarageCars       1460 non-null    int64
62 GarageArea       1460 non-null    int64
63 GarageQual       1379 non-null    object
64 GarageCond       1379 non-null    object
65 PavedDrive       1460 non-null    object
66 WoodDeckSF       1460 non-null    int64
67 OpenPorchSF      1460 non-null    int64
68 EnclosedPorch    1460 non-null    int64
69 3SsnPorch        1460 non-null    int64
70 ScreenPorch      1460 non-null    int64
71 PoolArea         1460 non-null    int64
72 PoolQC           7 non-null       object
73 Fence            281 non-null     object
74 MiscFeature       54 non-null      object
75 MiscVal          1460 non-null    int64
76 MoSold           1460 non-null    int64
77 YrSold            1460 non-null    int64
78 SaleType         1460 non-null    object
79 SaleCondition     1460 non-null    object
80 SalePrice_Log    1460 non-null    float64

```

dtypes: float64(4), int64(34), object(43)

memory usage: 924.0+ KB

None


In [29]:

```
df.head()
```

Out[29]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	Ui
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	/
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	/
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	/
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	/
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	/

5 rows × 81 columns



In [17]:

```
# Basic summary:  
df['SalePrice'].describe()
```

Out[17]:

```
count      1460.000000  
mean       180921.195890  
std        79442.502883  
min        34900.000000  
25%        129975.000000  
50%        163000.000000  
75%        214000.000000  
max        755000.000000  
Name: SalePrice, dtype: float64
```

In [18]:


```
sns.boxplot(df['SalePrice'])
```

C:\Users\Aarushi\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[18]:

```
<AxesSubplot:xlabel='SalePrice'>
```



There are two outliers with prices more than 700000.


**### The target variable : Distribution of SalePrice**

In [19]:

```
sns.displot(df['SalePrice']);  
#skewness and kurtosis  
print("Skewness: %f" % df['SalePrice'].skew())  
print("Kurtosis: %f" % df['SalePrice'].kurt())
```

Skewness: 1.882876

Kurtosis: 6.536282



As we see, the target variable SalePrice is not normally distributed. This can reduce the performance of the ML regression models because some assume normal distribution, see sklearn info on preprocessing. Therefore we make a log transformation, the resulting distribution looks much better.


In [20]:

```
df['SalePrice_Log'] = np.log(df['SalePrice'])

sns.displot(df['SalePrice_Log']);
# skewness and kurtosis
print("Skewness: %f" % df['SalePrice_Log'].skew())
print("Kurtosis: %f" % df['SalePrice_Log'].kurt())
# dropping old column
df.drop('SalePrice', axis=1, inplace=True)
```

Skewness: 0.121335

Kurtosis: 0.809532




In [8]:

```
# The Density Plot of SalePrice
sns.displot(df['SalePrice'])
```

Out[8]:

<seaborn.axisgrid.FacetGrid at 0x21011665520>



it is right-skewed distribution with the pick around 160k and quite long tail with maximum about 800k.

In [9]:

```
# Positive Skeweness:
df['SalePrice'].skew()
```

Out[9]:

1.8828757597682129

In [63]:

```
y_train = df["SalePrice"]
```

## Check Missing Data

Columns with Nan values



At first, I am checking the fraction of Nan values in each column.

In [67]:

```
total = df.isnull().sum().sort_values(ascending=False)
percent = (df.isnull().sum()/df.isnull().count()).sort_values(ascending=False)*100
missing_data = pd.concat([total, percent], axis=1, keys=['Total', 'Percent'])
missing_data.head(20)
```

Out[67]:

	Total	Percent
<b>PoolQC</b>	1453	99.520548
<b>MiscFeature</b>	1406	96.301370
<b>Alley</b>	1369	93.767123
<b>Fence</b>	1179	80.753425
<b>FireplaceQu</b>	690	47.260274
<b>LotFrontage</b>	259	17.739726
<b>GarageYrBlt</b>	81	5.547945
<b>GarageCond</b>	81	5.547945
<b>GarageType</b>	81	5.547945
<b>GarageFinish</b>	81	5.547945
<b>GarageQual</b>	81	5.547945
<b>BsmtFinType2</b>	38	2.602740
<b>BsmtExposure</b>	38	2.602740
<b>BsmtQual</b>	37	2.534247
<b>BsmtCond</b>	37	2.534247
<b>BsmtFinType1</b>	37	2.534247
<b>MasVnrArea</b>	8	0.547945
<b>MasVnrType</b>	8	0.547945
<b>Electrical</b>	1	0.068493
<b>Id</b>	0	0.000000


Further narrowing down, we will look for numerical type features with missing values in each of the train data and test data. First, we will search from the train data.

## Visualize with a heatmap

Let's visualize the correlation of feature quantities with a heatmap and confirm.

In [69]:

```
#correlation matrix
df=df.drop(['Id'],axis=1).copy()
corrmat = df.corr()
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corrmat, vmax=.8, square=True);
```



We can see that 'OverallQual', 'GrLivArea', 'GarageCars', 'GarageArea', 'TotalBsmtSF' etc are strongly correlated with 'SalePrice'.

## Analyze, identify patterns, and explore the data.

Analyze by pivoting features We will explore the relationship with 'SalePrice' about the feature quantity of the object type which does not include the missing value. From the name of feature quantity we pick up what seems to be strongly related to 'SalePrice' and analyze the mutual relationship with pivot.

## Creating new feature extracting from existing

We can convert the categorical titles to ordinal.

In [70]:

```
df[['HouseStyle', 'SalePrice']].groupby(['HouseStyle'], as_index=False).mean().sort_values(
```

Out[70]:


	HouseStyle	SalePrice
3	2.5Fin	2.579223
5	2Story	2.578496
7	SLvl	2.564995
2	1Story	2.563969
4	2.5Unf	2.557606
0	1.5Fin	2.550170
6	SFoyer	2.548212
1	1.5Unf	2.533208

'SalePrice' correlation matrix (zoomed heatmap style)

Select 10 features including 'SalePrice' which has strong correlation with 'SalePrice' and display it with a heat map.

In [71]:

```
#saleprice correlation matrix
k = 10 #number of variables for heatmap
corrmat = df.corr()
cols = corrmat.nlargest(k, 'SalePrice')['SalePrice'].index
cm = np.corrcoef(df[cols].values.T)
sns.set(font_scale=1.25)
f,ax=plt.subplots(figsize=(12,9))
hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kws={'size': 12},
plt.show()
```



'GarageArea' and 'GarageCars' are indicators showing the same thing at different angles, so we can see that the correlation coefficient is close.


Unfortunately, it was confirmed that 'HouseStyle' and 'HeatingQC' are weakly correlated with 'SalePrice'.

## **Scatter plots between 'SalePrice' and correlated variables**

In order to check for outliers, convert the heatmap to a scatter plot with the pairplot method.

In [74]:

```
#scatter plots
sns.set()
sns.pairplot(df[cols], height = 3)
plt.show();
```



## Bivariate analysis

When checking the correlation between 'SalePrice' and 'GrLivArea' with a scatter plots, outlier 2 is found at the lower right. Let's delete this.


Let's check the result of the operation with the scatter plot of 'SalePrice' and 'GrLivArea'.

In [82]:

```
# bivariate analysis saleprice/grlivarea
var = 'GrLivArea'
data1 = pd.concat([df['SalePrice'], df[var]], axis=1)
sns.scatterplot(data = data1, x=var, y='SalePrice')
```

Out[82]:

<AxesSubplot:xlabel='GrLivArea', ylabel='SalePrice'>



In this way the outliers have been properly deleted.

We will look at the heatmap of 'SalePrice' and feature amount best ten 10 again. The correlation coefficient also rises somewhat, which shows that the ranking has changed.

As we see, the target variable SalePrice is not normally distributed. This can reduce the performance of the ML regression models because some assume normal distribution, see sklearn info on preprocessing Therefore we make a log transformation, the resulting distribution looks much better.